



Upgrade to Ansible 12/ansible-core 2.19 and above

Ansible Style Guide

CfgMgmtCamp Ghent 2026
03.02.2026

Kostiantyn Volenbovskyi
Abacus Umantis AG



Agenda

Upgrade to Ansible 12 and above (ansible-core 2.19 and above)

Ansible Style Guide

- Variables
- Tasks
- (some) Modules
- Jinja, filters, plugins

Installation of Ansible

Option 1: Your Linux distribution is likely to have Ansible

- It is likely to be outdated (even after e.g. dnf update ☺)

Option 2: Installation via pip (sometimes pipx) in a Python virtual environment

- Should be analysed considering Python version on Ansible control node and especially on managed nodes

Option 2.1 Multiple virtual environments depending on Managed nodes?

Option 3. Ansible Execution Environments

Ansible versions overview (February 2026)

Ansible community has certain ansible-core version as dependency		Requirements			Notes
Ansible community package version	ansible-core package version	Control node Python	Managed node Python	Managed node PowerShell	
14	2.21	3.12 - 3.14 (* I guess)	3.9 - 3.14 (* I guess)	5.1 (I guess)	Not released, to be released in June 2026
13	2.20	3.12 - 3.14	3.9 - 3.14	5.1	
12	2.19	3.11 - 3.13	3.8 - 3.13	5.1	
11	2.18	3.11 - 3.13	3.8 - 3.13	5.1	EOL in December 2025
10	2.17	3.10 - 3.12	3.7 - 3.12	5.1	EOL
9	2.16	3.10 - 3.12	2.7 / 3.6 - 3.12	3 - 5.1	EOL Last version compatible with CentOS 7 and AlmaLinux 8, Debian 10 (and SLES12?)
8	2.15	3.9 - 3.11	2.7 / 3.5 - 3.11	3 - 5.1	EOL

In previous Ansible versions

- `| succeeded` is not supported anymore, use `is succeeded`
- `include:` is not supported anymore, use `include_role` (or `import_role`)
- `pre_tasks/post_tasks` in playbooks are not necessary anymore, just use tasks with mix of tasks and e.g. `import_role`
- `is truthy/is falsy` was implemented in Ansible 2.10

Ansible 12 upgrade/Ansible-core 2.19 upgrade

- 😊 As with any newer Ansible versions: compatibility with newer collections (newer modules/new modules...)
- 😊 Faster templating leads to faster execution of certain tasks
- 😊 Better error reporting
- 😢 Forces to implement conditionals in a certain way

Workaround that will work till ansible-core 2.23: `ALLOW_BROKEN_CONDITIONALS=true` in `ansible.cfg`

Check:

```
grep "Broken conditionals are currently allowed because" * -B 8 <directory with Ansible logs>*.log
```

- 😢 Evaluates variables in task attributes even in skipped tasks, eg. `delegate_to`.
- 😢 Evaluates variables in `import_tasks` that were skipped (?)
- 😢 Existing code must be changed
 - Deprecation is not «quick» when you update often but it is «quick» when you update Ansible rarely
 - 1-2 changes are needed per year in «the whole codebase». Sometimes via e.g. «sed» but not always...
- 😢 Certain improvements are not improvements but regressions ?

Ansible 12 upgrade: conditionals

No more support of conditionals like

(vip_ip_address is a string)

~~when: vip_ip_address~~

What needs to be done?

- Variables in conditionals (when:) should contain one of the following:
 - Strings: is truthy/is falsy; in this case when: vip_ip_address is truthy
 - Lists/dictionaries: | length > 0
 - Booleans or strings that can be converted to booleans: | bool.
- (it is not necessary for actual Booleans but it is recommended so that it less risk of confusion)
- ... or there is Ansible test («in», «is») : this results in a Boolean

Conditionals: you will need |bool

Below will result in a string, no matter spelling of true/false etc..

You should use |bool

```
input_verification_phase1: >-
{%- if trg_solution |d('') is truthy and trg_solution_id|d(0)|int is truthy -%}
true
{%- elif src_solution_id |d(0)|int is truthy -%}
true
{%- else -%}
false
{%- endif -%}
```

Ansible 12 upgrade: delegate_to

- Sometimes a «virtual host» is used
 - «virtual»: it is not in an inventory; nor it is added via `add_host`
 - it is sometimes part of a conditional
- What needs to be done:
 - It can't be an empty value ☹, it shouldn't be false
 - These variables in `delegate_to` should then have:
 - `| d` (omit)
 - Or «fake», non-empty string (for example "`void_host`")

Ansible 12 upgrade: undefined variables as values of task attributes

For example: value of `delegate_to` can never be undefined variable in all tasks (including skipped ones):

```
[ERROR]: Task failed: 'new_ca_renewalmaster' is undefined  
  
Task failed.  
Origin: /repo/roles/umantis.infra_ipaserver/tasks/changefreeipamaster.yml:1:3  
  
1 - name: Change CA renewal master  
    ^ column 3  
  
<<< caused by >>>  
  
'new_ca_renewalmaster' is undefined  
Origin: home/repo/roles/umantis.infra_ipaserver/tasks/changefreeipamaster.yml:3:16  
  
1 - name: Change CA renewal master  
2   shell: "ipa config-mod --ca-renewal-master-server {{ new_ca_renewalmaster }}"  
3   delegate_to: "{{ new_ca_renewalmaster }}"  
    ^ column 16  
  
fatal: [freeipahost.domain.com -> {{ new_ca_renewalmaster }}]: FAILED! =>  
  changed: false  
  msg: 'Task failed: ''new_ca_renewalmaster''' is undefined'
```

References

<https://forum.ansible.com/t/python-3-7-impact-on-el8-future-for-el9/6229>

https://docs.ansible.com/projects/ansible/latest/reference_appendices/release_and_maintenance.html#support-life

<https://docs.adfinis.com/ansible-guide/>



Ansible Style Guide

CfgMgmtCamp Ghent 2026
03.02.2026

Kostiantyn Volenbovskyi
Abacus Umantis AG



Acknowledgements/credits

Thanks to:

- My team: Site Reliability Team of Ansible Umantis
- Ansible core team
- Red Hat team that maintains Automation Good Practices;
- Companies that created their Style Guides: LinuxFabrik; Whitecloud
- Style Guide and other Ansible-related work: Thomas Jungbauer; Felix Fontein, Thomas C. Foulds, George Shuklin; Andreas Sommer; Vladimir Botka; Jeff Geerling...



HR Software for DACH (Germany, Austria, Switzerland)

- Applicant Management/Employee Management, On-/Off-boarding, Expense Management....
- Salary module

Leader in ERP- Software in Switzerland

- The largest independent SW vendor of Business Software for SME, founded in 1985
- Software for Finance, HR, Administration, Production, etc.
- 65'000 happy customers

Objectives

- To have an Ansible Style guide that:
 - leads to code that is written faster without compromises on quality - understood (results in less cognitive load)/maintained/debugged/reused easier...
 - doesn't suggest the code/code style that introduces a fault into managed node

also:

- Ansible code that has reasonable performance

(Somewhat) contentious questions

Contentious question	Alternative 1 (preferred by me)	Alternative 2
1. Dictionary notation	With dots (possible in 95% of cases)	With square brackets
2. Prefixing variables	Prefixes with role name so variable name is unique in the codebase	Additional Prefixes like: Internal : double underscore _ Registered variable: r_
3. Format of inventory	INI	YAML
4. Usage of quotes	Only if necessary	All strings are to be quoted
5. Quotes option	<ul style="list-style-type: none">- External quotes: ""- Internal quotes: ''- Additional quotes – use YAML multiline- >-	<ul style="list-style-type: none">- Single quotes- ''
6. Verb being used as task name	Verbs like: “Install” “Create” “Configure”	“Ensure”

Types of variables

- String
- Integer
- Floats (fairly rare)
- List
 - Typically you use element from the list, like [0], [1]
- Boolean: true and false (other options are possible though!)
- Dictionaries


mydictionary:

```
  name: rene  
  lastname: magritte
```

- List of dictionaries (it is a list but there are dictionaries)

mylistofdictionaries :

```
- name: rene  
  lastname: magritte  
- name: paul  
  lastname: van Haver  
  stagename: stromae
```



Handling of conditional/boolean values

VAR-7

- Instead of:

```
myapp_upgrade_needed | d(True) == True
```

use:

```
myapp_upgrade_needed | bool
```

Instead of:

```
myapp_upgrade_needed | d(false) == False
```

use

```
not upgrade_needed | d(true) | bool
```

Mixing strings and None: truthy/falsy

VAR-8

- Always use `else` in Jinja conditionals
- You might choose not to put anything after `else` : it will result in `None` and typically it should be fine (or you should implement e.g. assert)

```
peer_host: >-
{%- if other_datacenter |bool -%}
webproxy01.{{ peer_dns_zone }}
{%- else -%}
{%- endif -%}
```

| `bool` : only strings "1", "on", "yes", and "true" (case insensitive), or the numeral 1 return true. So you should use: when: `peer_host` is truthy


Variables: Jinja

Jinja2 allows to implement:

- Templates:
 - **MODULE-5**
 - templates should be preferred to 'files', they are much more future-proof

In templates or in variables:

- Cleaner conditionals (**if/elif/else**) for variables
- Loops using `for` (much faster than `loop:` and typically would give more features; but it is more complex)
- Changing values of variables



Conditionals

```
# Simpler, one-liner:  
vm_net_bridge_name: "{{ 'virtmgmt' if vm_net == 'infra_mgmt' and virt_node |bool else vm_net }}"  
  
vault_installation_required: "{{ true if installed_vault_version.stdout |d('') != vault_version else false }}"
```

```
# More complex: Jinja  
src_tenant: >-  
{%- if src_tenant_stable.rc == 0 and src_tenant_beta.rc == 1 -%}  
{{ src_tenant_stable }}  
{%- elif src_tenant_stable.rc == 1 and src_tenant_beta.rc == 0 -%}  
{{ src_tenant_beta }}  
{%- else -%}  
{%- endif -%}
```

Integers vs strings

- Integers:
 - implement comparison and arithmetic operations
- In case you don't need arithmetics:
 - You can't easily specify integer via extra-vars, so sometimes you use `| int` to cast to integer

Defining default values/ |default filter

VAR-9

- Default value for boolean: |d(false) (sometimes you will end up using |d(true))
- Default value for a string: |d('')
- Default value for list |d([])
- Default value for dictionary |d({})

List of dictionaries

- **VAR-10**
- Recommended structure of variables that are more complex: **list of dictionaries**
- It is fairly easy to implement many things in list of dictionaries (in configuration_properties there are keys with name value :
 - loop: just specify {{ configuration_properties }} then use {{ item.value }}
 - Retrieval of a value: selectattr
 - Create list of certain values: map
 - More complex: write Jinja

https://docs.ansible.com/ansible/latest/collections/ansible/builtin/subelements_lookup.html

https://docs.ansible.com/ansible/latest/playbook_guide/complex_data_manipulation.html

Nested loops/subelements

- **VAR-11**
- «One variable»: there is a need to have a **list inside a dictionary** that is inside of a list of dictionaries and iterate through that
 - Use: `| subelements filter`
- «Two variables»: there is a need to have **separate** list that should be executed for every element in the first list
 - For example there are 4 databases (`postgres_databases`) and for each one you need to enable 4 Postgres extensions (`postgres_extensions`)
 - Use: `with_nest`


Defining default values/ |default filter

- **VAR-12**
 - **Option 1 recommended** Try to initialize variables in defaults/main.yml
 - And then remove «is defined» statements; no more «variable not defined» errors
 - No need to use |d everywhere
 - **Option 2**
 - Use |d a lot

|d will **always** be used on variables coming from «register» statements

Variable definition 101

- defaults/main.yml
 - could be overriden by group_vars and/or extra_vars
 - vars/main.yml
 - Those are **constants**: they have priority over defaults/group_vars. Please do not override them via extra vars
 - Then there are:
 - group_vars (but remember about DRY and 'all' group)
 - Inventory variables
 - set_fact
- And other places – but do you need them?



vars vs. defaults; to_nice_yaml

- **VAR-13**
 - Variable used by a single role
 - Typically: `defaults/main.yml` or `vars/main.yml`;
 - Variable used by several roles:
 - Variable used by **two (or 3)** roles: `include_vars`
 - When it is used more than
 - Alternative 1: `group_vars/all/`: particular when roles have nothing to do with each other
 - Alternative 2: more `include_vars`
- `vars/main.yml` vs. `defaults/main.yml`
- In case it is overridden by `group_vars/extravars`: then it should be in `defaults/main.yml`
 - In case you consider that sometime in future this configuration might change because of SW version change: put it to `defaults/main.yml` and not in `vars/main.yml`

Additional files. In case, for example variables that contain different messages:

`vars/communication.yml`

It won't be »sourced» automatically, you will need `include_vars`

Use cases of `set_fact`

Typical use cases:

- «finalizing» the variable, eg. a password
- «propagating» variable to other hosts (`set_fact with delegate_facts`)
- it is done in a loop (but maybe you could use Jinja?), `include_tasks` implemented via a loop

Overwrite variable «dynamically»:

- in the beginning of the execution variable has certain value, but you need to get it changed (or you overwrite Ansible built-in variable)

A pitfall, static nature of `import_tasks`:

`set_fact` and `when`: you can't use `import_tasks`, you should use `include_tasks`

Recommendations

VAR-13

- A lot of `set_fact` usage can be avoided
 - typically to be replaced by `vars/main.yml`

VAR-14

- Avoid ‘variable undefined’ error by going through checklist:
 - Does a variable have a sensible default?
 - Use `defaults/main.yml` or other location
 - Empty value is sometimes fine but sometimes you need to avoid that (`«assert»`) gets into e.g. as script parameter unintentionally
 - Is the variable a «runtime» one (`register`)? Then use `|d` filter

Extra-vars vs. tags

- Extra-vars

VAR-15:

- Extra vars should not be mandatory for Day0/Day1 playbooks

VAR-16:

- Extra vars must be documented, typically should be verified by assert/fail
- Tags
 - Shorten the execution time/show clearly which tasks are executed/speed up troubleshooting
- **TASKS-6**
 - There should be no changes in outcomes of «execute with all tags» vs. «execute with a specific tag»

Debugging

When writing play, use tags and debug statements:

- Tags:
 - Some of the tags are useful in normal operation (less time for executing some exact task)
 - Tags are very useful in troubleshooting
 - Some of the tags are useful also in the maintenance of Ansible code
- Debug

Two main alternatives:

```
debug:
```

```
    msg: "{{ variable_to_debug }}"
```

- Typically leave them in code via `verbosity: 1`

```
pause:
```

```
    prompt: "{{ variable_to_debug }} "
```

Use `to_nice_yaml` filter in most of cases to make nicer output of various eg. dictionaries/lists within eg. «debug» module

Meta subdirectory

- In case you publish to Galaxy: OK!
- Otherwise look at your meta/main.yml:
 - Does anyone read/update it?
 - It is not used automatically by most of Ansible functionality
 - Do not use dependencies there, use import_role

TASKS2

meta/ directory can be removed completely in most cases

Blocks

TASKS-3

Main use case:


- You must use `block` when you would like to use `rescue`:

Other use cases:

- 2 or 3 tasks with common task attributes: maybe `block` is better, up to you!
- More than 3 tasks: consider using separate file and `import_tasks`

import_tasks/include_tasks

- **TASKS-4** Prefer import_tasks over include_tasks in case possible
- import_tasks is good because it can decrease the amount of same values for certain attributes over and over, attributes like:
 - become:
 - become_user:
 - run_once:
 - delegate_to:
 - tags:



Assert vs. fail

MODULE-6

- `ansible.builtin.assert/ansible.builtin.fail` modules implement «fail early» principle and provide user-friendly failure handling
- One of the main parameters is information that you give within: it is meant to be used by user executing Ansible
- Any playbook with more complicated input or somewhat «unreliable output» will benefit from assert/fail
- For multiple conditions `assert:` seems better than `fail`

Assert

MODULE-7

- `ansible.builtin.assert`

You might skip using `fail_msg` in case name of the task describes the condition

MODULE-3

- In case there are multiple conditions:
 - make sure that they are on separate lines and make sure that `fail_msg` contains text that covers both conditions!

Ansible filters vs. Jinja built-in methods

- Ansible filters
 - Pipe notation
 - Ansible tests
- Jinja built-in methods
 - Very similar to Ansible filters
- Dot notation

Lookup plugin(s)

- Always executed on Ansible control node (!)
- Functionality of various lookup plugins:
 - Generate random password
 - Check file contents
 - Find files that match the pattern (`find`)
 - Execute some Linux command (so that output of command is in Ansible variable)
 - For example `date` or `dig`

MODULE-4

```
# An example of the value for variable below: 2024-08-16T13:42:45.906+0200
```

```
current_date_full: "{{ lookup('pipe','date +%Y-%m-%dT%H:%M:%S.%3N%z') }}"
```

- Lookup value of variable so that it is passed by value instead of passed by reference

Comparison/search/grep

FILTER_TEST_JINJA-1

- Equal to == ; not equal to: != ; Comparison: > and < , typically with casting to **integer** (| int)
- Check if particular element is a element in a list or substring is part of a string:

in

- Search is done via Ansible **test**:

```
release_jira_description | lower is search('secondary')
```

- **regex_search/regex.findall**: Ansible filters that extract part of line

```
k8stool_download.dest | regex_search('.*\.(gz|zip|bz2|xz)$')  
found_rpms_base: "{{ found_issues.meta.issues[0].fields.description |  
regex.findall('application-(?:package1| package2)_(?:beta|stable).*\?\.\.rpm') }}"
```

Checking if file exists/read file

Ansible Control Node:

FILTER_TEST_JINJA-2

Check if file exists in Ansible control node:

```
hosts_path: /etc/hosts
```

...

```
debug:
```

```
    msg: "host file exists"
```

```
when: hosts_path is exists
```

Read file contents , file is on Ansible control node:

```
lookup/'file':  
  
info_new_dc: "{{ lookup('file', 'somedirectory/example.yml') | from_yaml |  
d([]) }}"
```

Managed Node:

MODULE-5

Checking if file exists in managed node

- ansible.builtin.stat
- exists dictionary key with boolean value

Read file contents, file is on managed node: slurp

References

RedHat Automation Good Practices:

https://redhat-cop.github.io/automation-good-practices/#_introduction

Linuxfabrik's Ansible Development Guidelines:

<https://github.com/Linuxfabrik/lfops/blob/main/CONTRIBUTING.rst>

Ansible Style Guide: Preparation Document, by [Thomas Jungbauer](#)

<https://github.com/tjungbauer/ansible-style-guide/blob/main/AnsibleStyleGuide.adoc>

<https://docs.adfinis.com/ansible-guide/>

25 Tips for Using Ansible in Large Projects by Thomas C.Foulds

<https://thomascfoulds.com/2021/09/29/25-tips-for-using-ansible-in-large-projects.html>

Adfinis Ansible Guide:

<https://docs.adfinis.com/ansible-guide/>

Ansible best practices by Andreas Sommer:

<https://andidog.de/blog/2017-04-24-ansible-best-practices>

References (2)

Ansible Best Practices by George Shuklin:

<https://www.slideshare.net/slideshow/best-practices-for-ansible/95475433#3>